

```
In [1]: import pandas as pd
```

```
In [2]: data=pd.read_csv("/home/placement/Desktop/BhanuSiva4K8/flat500.csv")
```

```
In [3]: data.describe()
```

```
Out[3]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [4]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   ID                    1538 non-null   int64  
1   model                 1538 non-null   object  
2   engine_power          1538 non-null   int64  
3   age_in_days           1538 non-null   int64  
4   km                    1538 non-null   int64  
5   previous_owners       1538 non-null   int64  
6   lat                   1538 non-null   float64 
7   lon                   1538 non-null   float64 
8   price                 1538 non-null   int64  
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [5]: `#df=data`  
`#data=df.loc[(df.model=='lounge')&(df.previous_owners==1)]`

In [6]: `#2-3`

In [7]: `data=data.drop(['ID','lat','lon'],axis=1)#unwanted columns removed`

```
In [8]: data
```

```
Out[8]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [9]: data=pd.get_dummies(data)
```

```
In [10]: data
```

```
Out[10]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [11]: data.shape
```

```
Out[11]: (1538, 8)
```

```
In [12]: #predicted value we removed from dataframe  
y=data['price']
```

In [13]:

y

```
Out[13]: 0      8900
          1      8800
          2      4200
          3      6000
          4      5700
          ...
        1533     5200
        1534     4600
        1535     7500
        1536     5990
        1537     7900
```

Name: price, Length: 1538, dtype: int64

In [14]: x=data.drop('price',axis=1)

In [15]:

x

```
Out[15]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...	...	...	...	...	...	...	...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

```
In [16]: #predicted value we removed from dataframe
y=data['price']
x=data.drop('price',axis=1)
```

```
In [17]: x
```

```
Out[17]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...	...	...	...	...	...	...	...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

```
In [18]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [19]: x_test.head(5)
```

```
Out[19]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0

```
In [20]: x_train.shape
```

```
Out[20]: (1030, 7)
```

```
In [21]: x_train.head(5)
```

```
Out[21]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0

```
In [22]: x_test
```

```
Out[22]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0
...	...	...	...	...	...	...	...
291	51	701	22000	1	1	0	0
596	51	3347	85500	1	0	1	0
1489	51	366	22148	1	0	1	0
1436	51	1797	61000	1	1	0	0
575	51	366	19112	1	1	0	0

508 rows × 7 columns

```
In [23]: y_train.head(5)
```

```
Out[23]: 527    9990
129     9500
602     7590
331     8750
323     9100
Name: price, dtype: int64
```



In [24]: `y_test`

Out[24]:

481	7900
76	7900
1502	9400
669	8500
1409	9700
...	
291	10900
596	5699
1489	9500
1436	6990
575	10900

Name: price, Length: 508, dtype: int64

In [25]: `x_train`

Out[25]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0
...	...	...	...	...	...	...	...
1130	51	1127	24000	1	1	0	0
1294	51	852	30000	1	1	0	0
860	51	3409	118000	1	0	1	0
1459	51	762	16700	1	1	0	0
1126	51	701	39207	1	1	0	0

1030 rows × 7 columns

```
In [26]: y_train
```

```
Out[26]: 527      9990
          129      9500
          602      7590
          331      8750
          323      9100
          ...
          1130     10990
          1294      9800
          860      5500
          1459      9990
          1126      8900
          Name: price, Length: 1030, dtype: int64
```

```
In [27]: from sklearn.linear_model import LinearRegression
          reg=LinearRegression()#creating object of LinearRegression
          reg.fit(x_train,y_train)#training and fitting LR object using training data
```

```
Out[27]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.  
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [28]: #x_test=[[1,51,1000,28800,3],[1,51,780,18800,1]]
```

```
In [29]: #above line to actual
```

```
In [30]: ypred=reg.predict(x_test)
```

In [31]: ypred

Out[31]: array([ 5867.6503378 , 7133.70142341, 9866.35776216, 9723.28874535,  
10039.59101162, 9654.07582608, 9673.14563045, 10118.70728123,  
9903.85952664, 9351.55828437, 10434.34963575, 7732.26255693,  
7698.67240131, 6565.95240435, 9662.90103518, 10373.20344286,  
9599.94844451, 7699.34400418, 4941.33017994, 10455.2719478 ,  
10370.51555682, 10391.60424404, 7529.06622456, 9952.37340054,  
7006.13845729, 9000.1780961 , 4798.36770637, 6953.10376491,  
7810.39767825, 9623.80497535, 7333.52158317, 5229.18705519,  
5398.21541073, 5157.65652129, 8948.63632836, 5666.62365159,  
9822.1231461 , 8258.46551788, 6279.2040404 , 8457.38443276,  
9773.86444066, 6767.04074749, 9182.99904787, 10210.05195479,  
8694.90545226, 10328.43369248, 9069.05761443, 8866.7826029 ,  
7058.39787506, 9073.33877162, 9412.68162121, 10293.69451263,  
10072.49011135, 6748.5794244 , 9785.95841801, 9354.09969973,  
9507.9444386 , 10443.01608254, 9795.31884316, 7197.84932877,  
10108.31707235, 7009.6597206 , 9853.90699412, 7146.87414965,  
6417.69133992, 9996.97382441, 9781.18795953, 8515.83255277,  
8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,  
8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,  
6020.70555100, 10025.02571520, 7270.77100000, 8411.45004000])

In [32]: `#savedmodel=pickle.load(open(flname,'rd'))`  
`#x_test=[[1,75,1062,8000,1]]`  
`#savedmodel.predicted(x_test)`

In [33]: `from sklearn.metrics import r2_score`  
`r2_score(y_test,ypred)`

Out[33]: 0.8415526986865394

In [34]: `from sklearn.metrics import mean_squared_error #calculating MSE`  
`mean_squared_error(ypred,y_test)`

Out[34]: 581887.727391353

```
In [36]: num=581887.727391353
print(num**0.5)
```

```
762.8156575420782
```

```
In [37]: y_test.head(5)
```

```
Out[37]: 481      7900
76      7900
1502    9400
669     8500
1409    9700
Name: price, dtype: int64
```

```
In [39]: ypred
```

```
Out[39]: array([ 5867.6503378 ,  7133.70142341,  9866.35776216,  9723.28874535,
10039.59101162,  9654.07582608,  9673.14563045, 10118.70728123,
  9903.85952664,  9351.55828437, 10434.34963575,  7732.26255693,
  7698.67240131,  6565.95240435,  9662.90103518, 10373.20344286,
  9599.94844451,  7699.34400418,  4941.33017994, 10455.2719478 ,
10370.51555682, 10391.60424404,  7529.06622456,  9952.37340054,
  7006.13845729,  9000.1780961 ,  4798.36770637,  6953.10376491,
  7810.39767825,  9623.80497535,  7333.52158317,  5229.18705519,
  5398.21541073,  5157.65652129,  8948.63632836,  5666.62365159,
  9822.1231461 ,  8258.46551788,  6279.2040404 ,  8457.38443276,
  9773.86444066,  6767.04074749,  9182.99904787, 10210.05195479,
  8694.90545226, 10328.43369248,  9069.05761443,  8866.7826029 ,
  7058.39787506,  9073.33877162,  9412.68162121, 10293.69451263,
10072.49011135,  6748.5794244 ,  9785.95841801,  9354.09969973,
  9507.9444386 , 10443.01608254,  9795.31884316,  7197.84932877,
10108.31707235,  7009.6597206 ,  9853.90699412,  7146.87414965,
  6417.69133992,  9996.97382441,  9781.18795953,  8515.83255277,
  8456.30006203,  6499.76668237,  7768.57829985,  6832.86406122,
  8347.96113362, 10439.02404036,  7356.43463051,  8562.56562053,
  8820.78555100, 10025.02571520,  7270.77100022,  8411.45004006])
```

```
In [ ]: #from sklearn.metrics import accuracy_score
#accuracy_score(y_test,ypred)
```

```
In [42]: #Results=pd.DataFrame(columns=['Actual',Predicted])
#Results['Actual']=y_test
Results=pd.DataFrame(columns=['Price', 'Predicted'])
Results['Price']=y_test
Results['Predicted']=ypred
#Result['km']=x_test['km']
Result=Results.reset_index()
Results['Id']=Results.index
Result.head(15)
```

Out[42]:

	index	Price	Predicted
0	481	7900	5867.650338
1	76	7900	7133.701423
2	1502	9400	9866.357762
3	669	8500	9723.288745
4	1409	9700	10039.591012
5	1414	9900	9654.075826
6	1089	9900	9673.145630
7	1507	9950	10118.707281
8	970	10700	9903.859527
9	1198	8999	9351.558284
10	1088	9890	10434.349636
11	576	7990	7732.262557
12	965	7380	7698.672401
13	1488	6800	6565.952404
14	1432	8900	9662.901035

```
In [44]: Results['DIFF']=Results.apply(lambda row:row.Price-row.Predicted,axis=1)
```

In [45]: Results

Out[45]:

	Price	Predicted	Id	DIFF
<b>481</b>	7900	5867.650338	481	2032.349662
<b>76</b>	7900	7133.701423	76	766.298577
<b>1502</b>	9400	9866.357762	1502	-466.357762
<b>669</b>	8500	9723.288745	669	-1223.288745
<b>1409</b>	9700	10039.591012	1409	-339.591012
...	...	...	...	...
<b>291</b>	10900	10032.665135	291	867.334865
<b>596</b>	5699	6281.536277	596	-582.536277
<b>1489</b>	9500	9986.327508	1489	-486.327508
<b>1436</b>	6990	8381.517020	1436	-1391.517020
<b>575</b>	10900	10371.142553	575	528.857447

508 rows × 4 columns

In [ ]: